

AI safety's institutional collapse creates a governance vacuum

The last two weeks mark the most concentrated period of AI safety erosion in the industry's history.

Between January 30 and February 13, 2026, at least 15 senior AI staff departed major labs (OpenAI, Anthropic, xAI), OpenAI dissolved its second safety team in two years and launched ads in ChatGPT, the U.S. government formally stripped "safety" from its AI standards body, and Palantir cemented itself as the de facto government AI partner—all while a 100-expert international report warned that voluntary AI governance is failing. For a new AI governance startup like CSOAI/EP³, this convergence of safety exodus, regulatory vacuum, and surging government demand represents a rare market-creation moment.

The resignation wave: 15+ departures in eight days

The week of February 6–13, 2026 saw an unprecedented cluster of senior AI departures spanning three major labs, driven by distinct but converging pressures.

At OpenAI, three events landed within 48 hours. **Zoë Hitzig**, a research scientist and Harvard junior fellow, resigned on February 10—the same day OpenAI began testing ads in ChatGPT—publishing a *New York Times* op-ed titled "OpenAI Is Making the Mistakes Facebook Made. I Quit." She warned that ChatGPT holds "an archive of human candor that has no precedent," including users' "medical fears, their relationship problems, their beliefs about God and the afterlife," (TechFlow) (SiliconANGLE) and that advertising built on this data "creates a potential for manipulating users in ways we don't have the tools to understand, let alone prevent."

(National Today) (CNN) The same week, reporting revealed that **Ryan Beiermeister**, VP of Product Policy, had been fired in early January after opposing OpenAI's planned "adult mode" for ChatGPT and raising concerns about child exploitation safeguards. OpenAI cited sexual discrimination against a male colleague; Beiermeister called the allegation "absolutely false." (Futurism) (CNN) Most consequentially, on February 11 OpenAI **disbanded its Mission Alignment team**—the six-to-seven-person group created in September 2024 to ensure AGI benefits humanity. Its leader, **Josh Achiam**, was reassigned to a new role as "chief futurist." (Baller Alert +2) This marks the **second dedicated safety team** OpenAI has dissolved; the Superalignment team was dismantled in May 2024 after co-leads Ilya Sutskever and Jan Leike departed. (FindArticles +2)

At Anthropic, the departure carried unique symbolic weight. **Mrinank Sharma**, who led the Safeguards Research Team since its inception in 2023, resigned on February 9 with a public letter (Worth) on X (eWEEK) (The Federal) that garnered nearly **one million views**. (The American Bazaar) (The Next Hint) Sharma—who holds a D.Phil. from Oxford in statistical machine learning (OpenTools.ai) and was affiliated with the Future of Humanity Institute—wrote: "The world is in peril. And not just from AI, or bioweapons, but from a whole series of interconnected crises unfolding in this very moment." (Worth +2) He directly critiqued Anthropic's internal governance: "Throughout my time here, I've repeatedly seen how hard it is to truly let our values govern our actions... where we constantly face pressures to set aside what matters most." (Futurism +2) His final research project documented thousands of incidents where AI assistants "distorted users' perception of reality on a daily basis." (UNILAD Tech) (The American Bazaar) Anthropic clarified that Sharma was "not the head of safety nor was he in charge of broader safeguards at the company," though multiple outlets had described him as such. (CNN) At least three other Anthropic staff—**Harsh Mehta**, **Behnam Neyshabur**, and **Dylan Scardinaro**—also departed around the same time. (PC Gamer) (eWEEK)

At **xAI**, the exodus was the largest by headcount. At least **nine public departures** occurred between February 6–11, including two co-founders: **Yuhuai (Tony) Wu** (reasoning lead) and **Jimmy Ba** (research/safety lead, who reported directly to Musk). Six of xAI's 12 original co-founders have now left. (CNN) Ba's farewell was notably ominous: "Recursive self improvement loops likely go live in the next 12 months... 2026 is gonna be insane and likely the busiest (and most consequential) year for the future of our species." (TechCrunch) Musk characterized the departures as a "reorganization" to speed growth; three departing staff announced plans to start new ventures together. (TechCrunch) The exits coincided with French authorities raiding X offices over Grok-generated nonconsensual deepfakes of women and children, and the pending SpaceX acquisition of xAI (TechCrunch) in a **~\$1.25 trillion all-stock deal**. (CNBC)

A critical pattern unites these departures: **no evidence of formal coordination** exists. As former OpenAI researcher Daniel Kokotajlo previously stated, "It's not been like a coordinated thing. I think it's just people sort of individually giving up." The timing cluster was driven by converging catalysts—OpenAI's ad rollout, Anthropic's pursuit of a reported **\$350 billion valuation**, (Business Today) and xAI's corporate restructuring—but the shared underlying concern is structural: commercial pressures are systematically overriding safety commitments at every major frontier lab. (Startup News) (CNN)

Palantir's rise as the government's AI gatekeeper

While safety researchers exit the private sector, Palantir Technologies has positioned itself as the indispensable bridge between Silicon Valley AI and U.S. government operations, achieving a level of institutional entrenchment with few parallels in defense technology history.

The numbers are staggering. Palantir's Q4 2025 earnings, reported February 2, showed **\$1.41 billion in revenue** (70% year-over-year growth—its highest ever as a public company), with U.S. government revenue hitting **\$570 million** (up 66%) (CNBC) and total contract value reaching a record **\$4.3 billion** for the quarter.

(Yahoo Finance) CEO Alex Karp called the results "indisputably the best results that I'm aware of in tech in the last decade." (CNBC) The company now holds a **\$10 billion, 10-year Army enterprise agreement** (Ainvest) (CCE | ONLINE NEWS) (the largest in its history), the **\$1.3 billion+ Maven Smart System** contract (Ainvest) for battlefield AI, (CCE | ONLINE NEWS) a **\$448 million Navy ShipOS** contract (The Motley Fool) (CCE | ONLINE NEWS) that reduced submarine schedule planning from 160 hours to 10 minutes, (The Motley Fool) (CCE | ONLINE NEWS) and dozens of other defense engagements. Its 2026 revenue guidance of **\$7.18–\$7.20 billion** implies 61% growth, (Tickeron) with U.S. commercial revenue expected to grow 115%. (Tickeron)

Palantir's policy influence matches its contractual footprint. The company submitted detailed recommendations to the White House OSTP's AI Action Plan process and subsequently declared it was "proud that the substance of so many of our recommendations helped shape the final plan." (OpenSecrets) Palantir is an anchor partner in the **U.S. Tech Force**, a program launched February 10 that recruits 1,000 private-sector technologists for two-year government stints (FedScoop) (Government Executive) at \$135,000–\$195,000 salaries, (Axios) with OPM Director Scott Kapor explicitly calling Palantir "no doubt a leader" in government AI. (Fox Business) A super PAC backed by Palantir allies and OpenAI/VC leaders, "**Leading the Future**," has pledged **more than \$100 million** to boost AI-friendly candidates in the 2026 midterms. (CNN)

Karp's public positioning mirrors the Trump administration's framing precisely. At Davos in late January, (Factually) he declared AI will "destroy humanities jobs" while boosting vocational workers, and warned Europe

is "falling behind." (Fox Business) At the Hudson Institute on February 13, he framed AI as a "binary battle" between the U.S. and China: "You can only explain the promise of AI if you understand and embrace the superiority of America and its culture. (Startup News) Second place is extinction." This rhetorical alignment with the administration's "**AI dominance**" doctrine is deliberate and strategic—it positions Palantir not just as a contractor but as an ideological partner.

The governance landscape has fractured along political lines

The regulatory environment for AI has shifted dramatically in the past two weeks, with safety infrastructure being actively dismantled at the federal level while state and international frameworks race to fill the gap.

The most consequential institutional change: the **U.S. AI Safety Institute was renamed the Center for AI Standards and Innovation (CAISI)**, with Commerce Secretary Howard Lutnick explicitly removing the word "safety" and declaring, "For far too long, censorship and regulations have been used under the guise of national security. Innovators will no longer be limited by these standards." (Technical.ly) CAISI's focus has narrowed to national security risks and global competitiveness. (National Institute of Standards ...) The AI Safety Institute Consortium (200+ organizations) remains nominally operational but is widely expected to be wound down. (Technical.ly) Meanwhile, NIST's budget faces a proposed **\$325 million cut** for FY2026, eliminating approximately 500 positions. (Government Executive)

The Pentagon's AI posture has shifted equally dramatically. Secretary of War Pete Hegseth's January 9 memorandum created an "AI-first" warfighting force (Inside Government Contracts) (Defense Mirror) with a **\$13.4 billion FY2026 budget**, (Ainvest) mandating "any lawful use" language in AI contracts and **explicitly dropping "responsible AI" and ethics language**. (Defense One) The GenAI.mil platform now offers Google Gemini and xAI's Grok across defense networks, (Defense Mirror) including classified systems. The Trump administration's December 2025 executive order established an AI Litigation Task Force to challenge state AI laws, (Kennedys Law LLP) (Benton Foundation) with the Commerce Department required to identify "onerous" state regulations (Kslaw) by March 11 (Kslaw) (Orrick) and the FTC directed to (Kslaw) classify state-mandated bias mitigation as a "per se deceptive trade practice." (Paul Hastings LLP)

State governments are pushing back hard. **New York's RAISE Act**, signed (Fisher Phillips) December 19, requires 72-hour incident reporting for frontier AI models (stricter than California's 15-day window), creates a new AI oversight office, (Fisher Phillips) and imposes penalties up to \$3 million per violation. (Skadden) **California's SB 53** took effect January 1, requiring frontier AI developers to publish safety protocols and report incidents. (CyberAdviser) **Colorado's AI Act** takes effect June 30 (Kslaw) (Seyfarth Shaw LLP) despite being called "probably the most excessive" state law in the Trump EO. (Kennedys Law LLP) **36 state attorneys general** have opposed any moratorium on state AI laws. (Orrick)

Internationally, the **EU AI Act** hit a major implementation milestone on February 2, with Commission guidelines on high-risk AI classification due (European Commission) and enforcement of Annex III obligations beginning August 2, 2026. (Pearl Cohen) The landmark **International AI Safety Report 2026**, published February 3 by 100+ experts led by Turing Award winner Yoshua Bengio, found that (Inside Global Tech) voluntary AI governance frameworks have doubled (12 companies published frontier safety frameworks in 2025) (Inside Privacy) but remain "insufficient"—creating an "evidence dilemma" for policymakers (Inside Global Tech) who lack data to verify corporate safety claims. The report explicitly recommends "defence-in-depth" layered

safeguards. (Internationalaisafetyreport) The **India AI Impact Summit** (February 16–20), the first global AI summit in the Global South, (Wikipedia +2) signals a deliberate shift in framing from "safety" to "impact." (iiss)

Anthropic and OpenAI have turned safety itself into a competitive weapon. Anthropic aired its **first Super Bowl ad** on February 8 directly mocking OpenAI's ChatGPT ad plans, (Worth +2) pledging Claude will remain "permanently ad-free," (Axios) and donated **\$20 million to Public First Action**, a political group backing state-level AI regulation—a direct counter to (Worth) the \$125 million "Leading the Future" PAC aligned with OpenAI. (Worth) Sam Altman called Anthropic's ads "clearly dishonest" (Fortune) and accused the company of being "authoritarian." (TechCrunch +2) This politicization of AI safety positions presents both opportunity and risk for new entrants.

What this means for a new AI governance venture

The convergence of safety exodus, regulatory vacuum, government demand, and corporate credibility crisis creates **five specific strategic openings** for CSOAI/EP³.

First, the "certified safe" verification gap is now acute. The International AI Safety Report explicitly identifies the inability to verify corporate safety claims as the central "evidence dilemma." (Inside Global Tech) With NIST's safety institute defanged, no independent body can credibly certify AI system safety. A new entity offering **third-party safety certification**—analogous to UL for electrical products or SOC 2 for cloud security—could fill a market need that both regulators and enterprises recognize. The EU AI Act's August 2026 enforcement deadline for high-risk systems creates immediate demand for compliance verification services, (Pearl Cohen) and state laws like New York's RAISE Act and California's SB 53 require incident reporting frameworks that companies will need help implementing.

Second, the departing safety researchers represent an unprecedented talent pool. At least 20+ senior safety researchers have left frontier labs in the past 18 months, (Gamebezz) many explicitly seeking new institutional homes. Mrinank Sharma's stated commitment to "courageous speech" (Worth) and human flourishing, (eWEEK) (The American Bazaar) Zoë Hitzig's expertise in AI governance and economic design, (National Today) (CNN) and the networks of former OpenAI Superalignment and Mission Alignment team members constitute a brain trust available for advisory boards, research partnerships, or direct hiring. Jimmy Ba's warning about "recursive self-improvement loops going live in the next 12 months" (TechCrunch) suggests the most technically credible voices believe the window for establishing governance infrastructure is extremely narrow.

Third, the Palantir/White House dynamic creates a government contracting pathway. Palantir's success demonstrates that the current administration values AI companies that frame capabilities in national security terms and offer practical deployment tools rather than abstract safety principles. CSOAI/EP³ could position not as a safety regulator (politically toxic in the current environment) but as an **operational assurance provider**—helping government agencies deploy AI with auditable decision pathways and verified performance claims. The U.S. Tech Force program, (FinTech Weekly) (Government Executive) NIST's CAISI (Government Executive) request for information on AI agent security (comments due March 9), (Federal Register) and the Pentagon's \$13.4 billion AI budget (Ainvest) represent concrete entry points. The key is framing governance as a capability enabler, not a constraint.

Fourth, the emotional intelligence / meta-governance layer addresses a documented market failure.

Sharma's final research documented thousands of daily incidents of AI systems distorting users' perception of reality. (UNILAD Tech) (The American Bazaar) Hitzig's resignation letter identified the manipulation potential of AI systems holding "an archive of human candor that has no precedent." (National Today) (CNN) Multiple wrongful death lawsuits allege ChatGPT reinforced suicidal ideation. Psychiatrists have documented "chatbot psychosis." (Tech Policy Press +2) California's Companion Chatbots Act (effective January 1, 2026) mandates safety protocols for harmful content and minor protections. (Kslaw) An EI³ framework that certifies AI systems' emotional safety and autonomy-preservation properties addresses a need that is both commercially viable (companies need compliance tools) and ethically compelling.

Fifth, the political polarization of AI governance creates space for a "neutral" institutional actor. With Anthropic and OpenAI waging a proxy war over regulation through competing super PACs (\$20 million vs. \$125 million), (Worth) and the Trump administration actively dismantling federal safety infrastructure while states build their own, the market lacks a credible, non-partisan governance institution. A new entity that works across political lines—helping companies comply with both federal "AI dominance" mandates and state safety requirements, bridging EU AI Act obligations with U.S. frameworks—could command significant market positioning.

Conclusion: the next 90 days are decisive

The structural picture is clear: **voluntary self-governance at frontier AI labs has failed under commercial pressure**, (Internationalaisafetyreport +3) federal safety infrastructure is being deliberately weakened, and the gap is being filled by a patchwork of state laws, (Kslaw) EU regulations, and Palantir-style government integration that prioritizes capability over safeguards. The safety researchers who understood these systems best are walking out the door. (Gamebezz)

For CSOAI/EI³, the actionable timeline is compressed. The NIST CAISI request for information on AI agent security closes **March 9**— (Federal Register) submitting substantive comments would establish institutional credibility. The EU AI Act's high-risk enforcement deadline of **August 2, 2026** (Pearl Cohen) creates a six-month window to build compliance certification services. New York's RAISE Act incident-reporting requirements (Fisher Phillips) and California's SB 53 safety protocol mandates generate immediate consulting demand. The India AI Impact Summit (February 16–20) (Wikipedia) offers a venue for positioning with the 30+ nations endorsing the International AI Safety Report. (Inside Global Tech) (Crowell & Moring)

The deepest strategic insight from this research is that the departing researchers' warnings are converging on a single thesis: **the institutions meant to govern AI from within have been captured by the commercial imperatives they were designed to constrain.** (Worth) (Creati.ai) The market opportunity for external, independent governance infrastructure has never been larger—or more urgent. The question is not whether this vacuum will be filled, but by whom, and whether they can move fast enough to matter before recursive self-improvement loops, as Jimmy Ba warns, "go live in the next 12 months."